

a Selection and reporting use different reconstruction anchors and image support

SELECTION



Source / latent



e15 candidates +
released full reference




Padded source



Tile M (e15) + R (released)

Reference: released / 444-MSE

REPORTING



Source / latent



e15 mixed + full



Crop to original extent







Final mixed reconstruction

Reference: e15 / 444-MSE or 6:1:1


CPU audit: 255 shared encoder/entropy tensors unchanged; 143/143 inherited synthesis tensors changed.

Required comparison: reconstruct both anchors, use the same valid pixels and metric, then verify the selected output.





a

A common delivered-loss cap


**b**

Router advantage after output verification






a A separate representation at each capacity



Each codec learns its own weights end to end. Matched-training D12 is needed to isolate depth.

b Select before encoding; transmit the selected expert






The decoder uses the signalled expert's entropy model and synthesis; expert latents are not interchangeable.




Charge selection, grouping, entropy resets, assembly and resident model memory to the complete system.

Prospective deployment · glyphs show architecture, not measured routing, throughput or batch support

Architectural capacity and explicit runtime assumptions

Exact parameter counts | Analytical runtime scenarios — not measurements or fitted forecasts





265 frame-QP pairs · equal frame weighting · 95% sequence-cluster intervals