


CPU REFERENCE: PADDING AND DEPTH SAVING

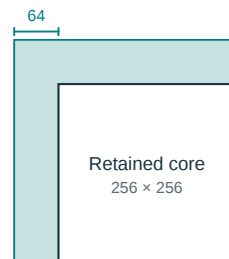
FUFREF1 image-pad64 geometry · bottom-right patch of a 2×2 grid

a 32-pixel halo · 288 → 320



Same coded area: 320 × 320

b 64-pixel halo · 320 → 320

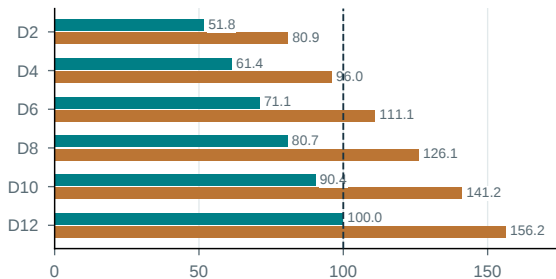


Same coded area: 320 × 320

Actual source context

Replicated padding

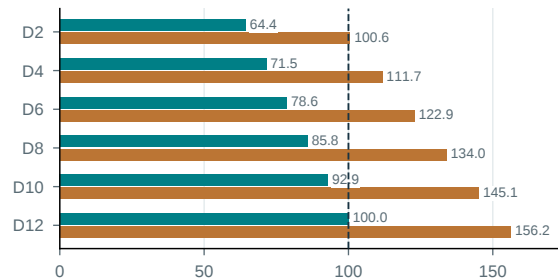
c Neural decoder



Conv2d MACs / full-frame D12 (%)

One 512 crop or four 256 cores

d Encoder with reconstruction

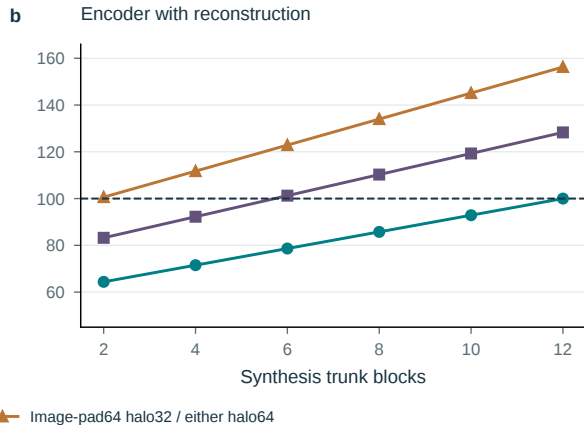
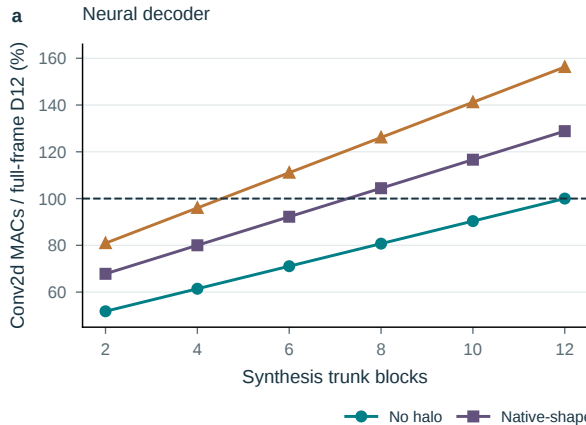


Conv2d MACs / full-frame D12 (%)

Four patches with halo 32 or 64

THE PADDING POLICY CHANGES THE COMPUTE ESTIMATE

Four 256 cores · native image-pad16 / latent-pad4 versus CPU reference image-pad64



Layer-by-layer Conv2d trace · native-shaped means geometry only · no CUDA timing, stream or numerical parity claim