

RegLIC: Region-Adaptive Learned Image Compression

Anonymous CVPR submission

Paper ID DRAFT

Abstract

*Learned image codecs usually apply one decoder depth to an entire image, although the benefit of additional synthesis varies across regions. We introduce **RegLIC**, a region-adaptive depth framework studied with the DCVC-UF intra codec. Its evaluated shared-exit path recovers one latent, executes later synthesis blocks only on surviving tiles, and assembles mixed-depth features through adapters and a common reconstruction head. Across 53 intra frames and five quality indices, learned routing saves 27.97% of modelled synthesis MACs at a nominal 0.1 dB YCbCr-MSE-ratio target under per-source calibration; 2.54 percentage points exceed ordered dithering, which mixes depths without content-dependent placement. A retrospective common delivered-loss cap preserves a routing margin, whereas a fixed-control QP32 replay does not yet resolve its paired advantage. We also define a complementary 2/4/6/8/10/12-block codec bank to study independent representation specialisation; its router is not part of the reported results. These controls distinguish savings from shallower decoding, spatial allocation, and the still-unmeasured effect on full-codec latency.*

1. Introduction

A learned codec normally spends the same synthesis computation on every part of an image. Yet another pair of decoder blocks may recover visible structure in one region and change almost nothing in its neighbour. Making each block faster lowers the cost of both regions; deciding *where the blocks run* addresses a different inefficiency. This is the question behind **RegLIC**: can a codec allocate reconstruction capacity locally without giving up the compression benefits of a shared representation?

Spatial capacity selection is established in restoration: ClassSR assigns super-resolution patches to networks of different sizes [6], while APE lets patches leave a shared network according to predicted incremental benefit [12]. AdaNIC brings spatially dynamic transforms to image compression [11]. For a codec, however, decoder depth is coupled to the coded representation and the rate at which it was



Figure 1. **One image, unequal decoder depth.** A conceptual cutaway shows regions sharing one coded representation while using different amounts of synthesis. This is an illustration, not a measured allocation or reconstruction.

produced. A visually complex patch need not benefit from a deeper decoder if its bitstream carries little recoverable detail. The relevant quantity is therefore the *marginal reconstruction gain* of more synthesis, relative to its computation and coding cost.

We study this allocation in the intra model of DCVC-UF [7]. Its twelve constant-width synthesis blocks expose depth as a clean capacity axis without changing feature width. RegLIC has two complementary realisations. In the *shared-exit* path, every region uses the same coded latent and decoder prefix, then exits after a chosen block pair. In the *independent-depth* path, a selector chooses a complete 2/4/6/8/10/12-block codec before encoding a region, allowing its analysis and entropy models to specialise as well. The latter trades representation reuse for specialised coding and introduces region-boundary and signalling costs. We evaluate the shared-exit path here; the independently trained bank is an ongoing capacity study, not a demonstrated routed system.

The evaluated design must actually *skip* computation:

computing every depth and choosing an output afterwards would save no decoding work. A full-frame stem supplies a common latent feature canvas; conditional block pairs process a shrinking batch of active tiles. Exit adapters align features from different depths before a shared full-frame repair and reconstruction head. This preserves a common output path while allowing each tile to stop at a different depth.

Separating the benefit of shallow decoding from the benefit of *placing* deeper computation is central to the evaluation. Uniform exits and ordered dithering already reduce average depth; the latter mixes depths without predicting which locations need them. At a nominal 0.1 dB YCbCr-MSE-ratio target across 53 intra frames and five QPs, spatial early exit saves 27.97% of modelled synthesis MACs under per-source calibration. The value of learned placement beyond simply using less depth is smaller and depends on the quality budget. When controls are fitted on other sequences, its QP32 paired advantage remains unresolved. The arithmetic saving is thus established more strongly than a source-independent routing gain or a complete-codec speedup.

Our contributions are a conditional DCVC-UF synthesis path with executed tile exits; a controlled comparison that isolates spatial placement from cheaper average depth; and a depth-bank protocol that tests the separate question of whether independent representation specialisation can repay its coding and boundary costs. Together they turn synthesis depth from a fixed architectural choice into an explicit, measurable spatial allocation problem.

2. Related Work

Making learned decoding cheaper. Shallow synthesis reduces decoder computation, with stronger analysis or iterative encoding compensating for lost capacity [15]. DCVC-RT emphasises memory traffic and function calls [4]; DCVC-UF combines chunk-based video coding with streamlined entropy interaction [7]. We modify its intra synthesis transform. A fixed shallow decoder controls for adaptation, while complete execution measurements distinguish arithmetic reduction from acceleration.

Choosing a network or choosing an exit. ClassSR learns to assign super-resolution patches to networks of different capacity [6]. APE instead lets patches leave a multi-exit network according to predicted incremental reconstruction benefit [12]; it already shares the reconstruction tail across exits. AdaRevD predicts patch difficulty to choose a deblurring sub-decoder [8]; RBQE terminates compressed-image enhancement early [13]. Swift introduces compute-dependent early exits in a layered neural video decoder, with different output resolutions [2]. RegLIC retains one entropy-decoded latent and fixed output resolution: adapters align conditional feature tiles before full-frame reconstruction. The proposed

independent-depth bank also selects the analysis transform and entropy model.

Adaptive capacity inside a codec. Slimmable compressive autoencoders jointly optimise rate and distortion at multiple network widths [14]. ELFIC allocates decoding complexity per instance [17]; ABC adapts backbone and autoregressive structure to content and computational budgets [16]. AdaNIC varies transform width spatially through blockwise operators and a router [11]. For video, cgSlimDecoder allocates decoder channels under complexity constraints [3]. MixCompress routes among different-depth experts using only the rate embedding, allowing assignments to be cached independently of image content [10]. RegLIC varies *depth* through nested, constant-width synthesis blocks and executes later stages on a shrinking tile batch. We measure the spatial placement premium over blind mixtures that already reduce average depth.

Spatial selection among complete codecs. Blard et al. train content-specialised codecs of the same architecture, search their regional rate-distortion costs, and signal a mode map [1]. They also process connected same-mode regions jointly; our merging diagnostic tests this existing strategy in DCVC-UF. The planned bank varies synthesis depth and investigates selection before candidate encoding. Shared exit retains one representation, separating representation reuse from independent-expert specialisation.

Preserving context across blocks. Kim et al. derive overlap and boundary rules that preserve full-image CNN reconstruction [5]. RegLIC instead accepts context loss in its zero-halo suffix and learns a common repair/head. Its tiled all-deep output must therefore be evaluated separately from full-frame synthesis.

3. Spatially Adaptive DCVC-UF

3.1. The allocation problem

Let $\mathbf{a} = (a_1, \dots, a_N)$ assign a synthesis option to each region. We seek the least computation that keeps the assembled reconstruction within a quality-loss budget:

$$\min_{\mathbf{a}} \mathcal{C}(\mathbf{a}) \quad \text{subject to} \quad \Delta(x, \hat{x}_{\mathbf{a}}; \hat{x}_{\text{ref}}) \leq \delta, \quad (1)$$

where $\mathcal{C}$ is the chosen compute cost and $\hat{x}_{\text{ref}}$ fixes the reconstruction anchor. The main sweep uses $\Delta_{444} = 10 \log_{10}[D_{444}(x, \hat{x}_{\mathbf{a}})/D_{444}(x, \hat{x}_{\text{ref}})]$, with $D_{444} = \|x - \hat{x}\|_2^2 / (3HW)$ on equal-channel, unclipped YCbCr 4:4:4 tensors. Positive values denote a quality loss in dB. The constraint belongs to the *assembled image*: spatial interactions can invalidate an additive tile-error approximation.

For early exit, a_t identifies a stopping point after one latent has been recovered. For the prospective model bank, a_t identifies a complete codec before encoding, making both payload rate and reconstruction depend on the assignment. The bank must therefore also be compared at matched rate, including expert signalling (Sec. 5). We implement shared exits below.

3.2. Shared latent, variable depth

Let $\hat{y}$ be the recovered latent of the fixed DCVC-UF analysis transform and entropy model. The synthesis network upsamples it by two and processes 384-channel features at one eighth of the image resolution. Four of its twelve trunk blocks form the full-frame stem; the remaining eight form four conditional block pairs.

The stem output is partitioned into 32×32 feature tiles, each covering 256×256 image pixels. For a tile assigned depth k , let $f_t^{(k)}$ be its feature after the corresponding block pair and let A_k be its adapter. The complete image is reconstructed as

$$\hat{x}_{\mathbf{k}} = H\left(q \odot G\left[\text{stitch}_t A_{k_t}(f_t^{(k_t)})\right]\right), \quad (2)$$

where G is grid-conditioned repair, q is the decoder's quality scaling, and H is the common output head, including pixel shuffle. Adaptation acts on features before the shared head, so tiles with different stopping depths need not be stitched directly in image space.

Reachable depth and actual execution. The reachable depths are $\mathcal{K} = \{6, 8, 10, 12\}$; depths two and four are inside the shared stem and are not evaluated exits. At each conditional stage the implementation gathers only the surviving tiles, executes the next block pair, and writes departing tiles through their adapters. A deeper tile reuses all computation already performed. Evaluating all four candidates and masking their reconstructions would not implement this execution saving. For the recorded router map at 0.1 dB, the active population decreases from 40 to 27, seven and three tiles (Fig. 3).

3.3. Exit alignment and boundary context

An intermediate feature is not automatically compatible with the final head. The exit adapter supplies the missing transformation with pointwise residual operators. Its residual output is zero-initialised, preserving the inherited feature at initialisation. Adapter capacity scales with the depth replaced. Pointwise operators avoid adding an exit-specific spatial receptive field; their arithmetic is included in the cost model. The deepest exit uses an identity connection rather than an additional learned adapter.

Independent tile execution still removes context from the trunk's spatial operators. The conditional trunk uses replicate

padding without a halo. After stitching, a depthwise spatial filter and pointwise mixing act on the full feature canvas. A learned gate indexed by position within the tile gives the repair access to the known boundary lattice. The head then operates on the complete canvas. This design places context recovery at the output stage rather than expanding every conditional block with an overlapping halo. We measure the trained checkpoint's dependence on adapters and repair in Sec. 4.7; matched retraining is needed to choose their capacity.

3.4. Learning exits and their selector

The evaluated shared-latent checkpoint, denoted e15, is initialised from released DCVC-UF weights. Analysis and entropy parameters remain frozen; the synthesis network, adapters and repair are fine-tuned. Training samples tile depths and applies reconstruction supervision through the mixed tiled path. Thus the adapters learn under the same stitching and full-frame head used at inference. The inspected training configuration also enables a full-depth reconstruction anchor against the frozen release and adjacent-exit feature distillation. Its objective can be written as

$$\mathcal{L} = \mathbb{E}_{\mathbf{k}}[R + \lambda D(x, \hat{x}_{\mathbf{k}})] + a\mathcal{L}_{\text{anchor}} + b\mathcal{L}_{\text{feature}}. \quad (3)$$

Here the latent-rate term R is shared across exits. With the front end frozen it does not optimise exit-dependent rate. The archived fine-tuning configuration uses 512-pixel crops and batch size eight; it is distinct from the new full 105-epoch official-recipe depth study in Sec. 5. The supplement records the available training provenance and its limits.

The stopping predictor projects the stem from 384 to 48 channels, pools each tile's mean and standard deviation, and incorporates normalised QP. An MLP with two 256-unit hidden layers produces exit log scores. For candidate tile errors $M_{t,k}$ and normalised costs c_k , its soft training labels favour low values of $M_{t,k} + \lambda_c c_k$. Soft-label cross-entropy and a regret term penalise costly decisions while allowing nearly equivalent exits to receive similar scores. The checkpoint uses $\lambda_c = 1.3 \times 10^{-5}$ and temperature ratio 0.1; the exact target, input layout and optimisation settings are given in the supplement. This objective learns a preference among exits, not a bound on final image error.

3.5. Stopping under a quality target

For equal padded tiles, the arithmetic model is

$$C(\mathbf{k}) = \frac{1}{N} \sum_t c_{k_t}, \quad S(\mathbf{k}) = 100[1 - C(\mathbf{k})]. \quad (4)$$

Shared stem and output work, exit adapters and conditional depth are included in c_k . Entropy-recovery networks, arithmetic coding, router execution, signalling and system overhead are outside this synthesis MAC quantity. The four costs





Figure 2. **Share the representation; spend further depth only on surviving tiles.** A four-block shared stem supplies the MLP with features and QP. Cost-adjusted scores select exits at total depths 6/8/10/12. Each departing feature passes through its adapter (identity at depth 12) into one canvas, followed by full-frame repair and reconstruction. The illustration follows the implemented connectivity; tensor contents and the output scene are conceptual, AI-generated artwork rather than measured reconstructions.

Spatial work disappears as tiles exit

Recorded 5×8 tile grid · one shared latent



77 / 160 tile-block-pair evaluations

Counts describe execution; they are not timings.

Figure 3. **A spatial decision becomes a smaller workload.** The recorded videoSRC05 map at QP32 leaves fewer tiles active at each successive block pair. Coloured cells run the indicated stage at their padded spatial coordinates; pale cells have already exited. Departures feed the aligned canvas in Fig. 2. The suffix evaluates 77 tile pairs versus 160 for an all-deep map; this work count does not measure GPU time.

are approximately $\{0.609, 0.758, 0.870, 1.010\}$ relative to full-frame synthesis. In particular, a tiled all-deep map can cost more than the full-frame reference because repair is still executed.

Source-informed search minimises $M_{t,k} + \lambda c_k$ per tile and adjusts λ for a target table loss. The learned alternative uses

$$k_t(\beta) = \arg \max_k \{\ell_{t,k} - \beta c_k\}, \quad (5)$$

where ℓ denotes log scores. Uniform depth selects one exit

per frame. Bayer dithering mixes adjacent exits with a fixed spatial ordering. Its mixture fraction is calibrated under the same target, so it controls for finer average depth without predicting spatial difficulty.

In the recorded sweep, all controls are calibrated per source frame using stored error tables. This measures the allocation available with encoder-side source access. The fixed-control experiment in Sec. 4.6 removes per-source calibration. Increasing β cannot increase selected cost, but source distortion need not be monotone; the supplement audits the finite assignment path for both router and dither.

3.6. Quality references and decision information

The table target compares padded YCbCr 4:4:4 candidate errors to a released-weight full-depth reference. Final reported loss instead compares the cropped mixed output to the cropped, fine-tuned e15 full-frame output. These anchors are not interchangeable: the encoder/entropy weights match, while inherited synthesis weights have changed. Full-frame repair also couples neighbouring tiles. A table threshold therefore does not guarantee the same numerical loss after final reconstruction. We retain this distinction in the nominal sweep and add a retrospective common-reference delivered-loss comparison. A separate QP32 replay on all 53 frames measures the e15 full-frame RGB loss against released synthesis as 0.00908 dB; details are in the supplement.

An encoder-selected map must be signalled to the decoder; alternatively, a decoder-local predictor must recover the same decisions from available features and a reproducible control. We use explicit two-bit exit indices in a CPU stream-only check (Sec. 4.7). Its common payload confirms the representation-sharing path, while headers and map bytes remain part of the complete rate. Source access is legitimate

at the encoder; its candidate evaluations and calibration must be charged to the decision cost.

4. Experiments

4.1. Evaluation protocol

We evaluate the first intra frame of 53 sequences at QPs $\{0, 16, 32, 48, 63\}$: 265 frame-QP pairs, six targets from 0.05 to 0.5 dB, and identical e15 weights, tile geometry and exit costs for all policies. The archive contains actual mixed reconstructions. Additional CPU replays test fixed controls, reference drift and component dependence. All decoder comparisons concern DCVC-UF; complete GPU latency is unmeasured.

Two complementary protocols answer different questions. The *nominal sweep* compares maps selected under the same table target, retaining the common feasible cases of all four rules: 204 pairs at 0.05 dB, 263 at 0.1 dB and 265 at each larger target. The *delivered-cap analysis* reselects from the outputs already measured at those targets under a common cropped YCbCr 4:4:4-loss cap, with the full-frame e15 output as an explicit fallback. It retains all 265 pairs at every cap. The second protocol is retrospective and source-aware; its cost omits acquiring and rejecting candidates.

Confidence intervals use 5,000 paired bootstrap draws over sequences, keeping QPs from the same sequence together. They quantify sequence-sample variation conditional on fixed weights and controls, without including training or calibration uncertainty. The supplement also reports a conservative content-group sensitivity check. A separate fixed-cohort comparison retains the same 204 pairs across the entire nominal sweep, avoiding changes of sample membership as the target relaxes.

4.2. The reconstruction value of depth

The stored uniform-exit reconstructions provide a direct depth control. Across all 265 pairs, their mean padded YCbCr 4:4:4 losses against the released full-frame reference are 0.205/0.097/0.061/0.038 dB at 6/8/10/12 blocks. Their modelled MAC savings are 39.12/24.22/13.03/−0.95%, respectively. Additional depth has diminishing average benefit and non-uniform local benefit: mean within-frame tile-gain IQRs are 0.075/0.034/0.024 dB for successive block pairs. The supplement gives the distributions and paired intervals. These measurements concern actual uniform maps on padded YCbCr 4:4:4; they motivate spatial allocation but are not cropped mixed-map losses or results for independently trained shallow codecs.

4.3. Early exit reduces work; routing improves its allocation

At the nominal 0.1 dB target, RegLIC with learned routing saves 27.97% of modelled synthesis MACs. Source-

informed search reaches 29.52%, while uniform depth and dithering reach 21.07% and 25.43%, respectively (Fig. 4). The ordering separates three effects. Shallower synthesis supplies most of the reduction. Mixing adjacent depths improves the available cost resolution. Content-dependent placement contributes a further 2.54 percentage points over dithering, with a paired 95% interval of [1.83,3.26].

4.4. The gain persists under a common delivered-loss cap

Does the margin persist when every selected output meets the same measured loss cap? For each policy and frame-QP pair, we deduplicate the maps measured at the six nominal targets, retain those meeting the cropped Δ_{444} cap within 10^{-4} dB, and choose the cheapest. Repeated maps have identical recorded losses and costs. The evaluated e15 full-frame output supplies a zero-loss, zero-saving fallback; no quality point is interpolated.

Figure 5 reports the 0.1 dB case. The router saves 27.75% versus 24.81% for dithering, a 2.93-point paired difference [2.11,3.79]. Their mean achieved losses are also close: 0.0786 and 0.0791 dB. Only one router case uses the full-frame fallback, compared with two dither cases. The margin survives a common measured cap. Selection remains retrospective and source-aware, omits candidate-acquisition cost, and searches only the recorded outputs.

4.5. Why the routing margin contracts

The cheapest reachable depth imposes a 39.12% ceiling on the arithmetic model. As the target relaxes, more tiles can take that exit and there are fewer consequential spatial decisions. The nominal router-dither margin declines from 2.54 points at 0.1 dB to 1.31 at 0.2 dB, 0.38 at 0.3 dB and 0.02 at 0.5 dB. Under the delivered-cap protocol the margin similarly falls from 2.93 to 0.35 points between 0.1 and 0.3 dB. The fixed-cohort result and QP breakdown in the supplement indicate that this behaviour is not explained solely by a changing feasible subset. At the 0.1 dB delivered cap, routing saves more in 170 pairs, less in 24, and ties in 71. At 0.5 dB, 262 of 265 pairs tie; the residual gain comes from one sequence.

The eight sources in Fig. 6 show how routing changes with content and budget. Relaxing the target admits more shallow exits, while routing, dithering and source search spend depth differently. A fixed sequence-name hash and dataset quotas select sources without using outcomes; infeasible cases remain. Luma supplies spatial context, not quality evidence.

A reconstruction intervention isolates placement from exit proportions. For every QP32 Q90 map, we replay three permutations within equal valid-tile extents, preserving depth counts and depth-weighted valid area: 53 frames and two policies give 424 map cases. Relative



a Budget shapes the compute frontier**b** Increment over dithering

Figure 4. **Early exit supplies the saving; spatial routing adds a budget-dependent margin.** **a**, Measured operating points shown in orthographic 3D; the policy axis is categorical and ribbons only guide the eye. **b**, Paired advantages over Bayer dithering, with 95% sequence-cluster intervals; the shaded row highlights 0.1 dB. QP and fixed-cohort analyses are in the supplement. All curves use per-source calibration. MAC savings exclude router and signalling overhead and are not measured runtime reductions. Targets use the Δ_{444} MSE-ratio loss defined in Sec. 3.




Figure 5. **Placement matters most before the cheap exit saturates.** **a**, Achieved mean loss and MAC saving for all four policies at the 0.1 dB cap. **b**, Paired router-dither saving differences across measured YCbCr 4:4:4-loss caps, with 95% sequence-cluster intervals. All 265 pairs remain; uniform/dither/router/search use the e15 full-frame fallback 11/2/1/2 times. Selection is retrospective and source-aware.

tations and the fixed development cohort. The intervention supports spatial value, not a net-runtime or external-test advantage.

4.6. Fixed controls test the routing advantage

To remove per-source calibration, we fit each scalar control on four sequence folds and apply it to the fifth, keeping all QPs together. Calibration minimises compute under a mean or 90th-percentile (Q90) table-loss target. Actual QP32 replay gives 318 reconstructions: 53 frames, three policies and two controls, with no quality fallback.

Figure 7 reports cropped RGB loss against e15 full-frame synthesis. Q90 calibration lowers the observed frequency of losses above 0.1 dB, at the cost of less arithmetic saving. The router-dither saving difference is 1.87 points under mean control, with a paired 95% interval of $[-1.60, 5.19]$, and 2.54 points under Q90, with interval $[-0.42, 5.54]$. Both include zero, and the policies do not have exactly matched achieved quality. Thus the source-calibrated margin has not yet been established for fixed controls.

This is a calibration diagnostic on development-used data, not an untouched test set. Two similar-content pairs cross the sequence folds; grouping them in a bootstrap sensitivity check leaves both intervals crossing zero. An aligned released-weight replay also measures the reference drift: e15 full-frame loses 0.00908 dB mean RGB PSNR against released D12. Using released D12 raises absolute policy losses by that same per-image offset without changing their paired differences. The supplement gives the split, predictor and metric audits.

to the mean shuffled RGB MSE, original router placement gains 0.0139 dB $[0.0095, 0.0193]$; dithering's gain is 0.0004 dB $[-0.0003, 0.0010]$. Unchanged maps remain included. These intervals condition on three sampled permu-




Figure 6. **Adaptive computation across content and budget.** a–h, Eight sources selected by sequence-name hash within fixed dataset quotas. Each row shows source luma, router maps at nominal 0.05/0.10/0.30 dB, and dithering/source-search maps at 0.10 dB. Colours encode executed depth on valid image extent; N/A retains infeasible cases. All maps are recorded at QP32 with source-calibrated controls. Source thumbnails show content, not reconstruction quality.




Figure 7. **Fixed controls trade saving against quality risk.** Actual reconstructions on all 53 QP32 frames. Colour identifies the policy; filled circles and open squares denote mean and Q90 calibration. Labels count RGB losses above 0.1 dB. Axes report achieved RGB loss against e15 full-frame and modelled synthesis MAC saving. Controls fitted on other sequence folds do not match achieved quality.

4.7. Checking the mechanism

Exit alignment and repair. With the checkpoint and 53 Q90 maps fixed, we replace adapters, repair, or both by identity and retain all images. Adapter removal costs 1.2309 dB mean RGB PSNR, with a paired 95% interval of [0.8765, 1.6582]. Repair removal costs only 0.00208 dB [0.00060, 0.00363]; removing both costs 1.2538 dB. The

trained model depends strongly on exit alignment. The small repair benefit motivates a matched-retraining and timing comparison before increasing its capacity. These interventions measure dependence of fixed weights, not the best performance attainable without each module.

Stream-only reconstruction. A separate CPU decoder receives the common entropy payload and explicit two-bit exit indices. On two predeclared QP32 frames with four maps each, all eight images, latents, symbols and entropy indexes match encoder-side reconstruction exactly; inner entropy bytes stay identical across maps. This verifies the shared-payload path on those cases. Map/header bytes and malformed-stream checks are in the supplement; native CUDA interoperability and full-cohort timing remain unmeasured.

4.8. What the encoder pays for a stopping decision

Source-informed search needs candidate reconstructions to populate its error table. The archived QP32 benchmark records 1,023.26 ms for four distinct candidate outputs and 336.87 ms for a shared tiled pass. These upper medians pool 24 trials over eight frames; raw timing distributions and device load were not retained. They indicate the cost of acquiring decision information, not an end-to-end speedup. The supplement details the stages.

A fixed-control predictor can avoid the error table; the source-calibrated router cannot, because its per-frame β uses the same information. Latency comparisons must charge feature generation, calibration, signalling, conditional execution and verification. Where learned and blind maps

459 nearly coincide, even small prediction overhead can erase
460 the benefit.

461 5. From Shared Exits to Independent Depths

462 Shared early exit commits every region to one encoded rep-
463 resentation. An independent codec can specialise its repre-
464 sentation for a smaller synthesis network, but forfeits that
465 reuse. This motivates our complementary 2/4/6/8/10/12-
466 block DCVC-UF bank, where a later ClassSR-inspired se-
467 lector will choose a codec before encoding each region.
468 Only pinned D2/D4/D6 interim milestones are analysed here;
469 D8/D10 and final matched models remain outside this com-
470 parison. D12 is the released anchor. No six-expert routing
471 result is claimed.

472 **A controlled capacity frontier.** The scratch runs pre-
473 serve the pinned official 105-epoch recipe [9], including
474 all 384,795 Open Images train_0/1/2 images, batch size 16
475 and FP32. Analysis and entropy weights learn alongside
476 synthesis, so equal QP does not imply equal coded rate. Two
477 fixed milestones, epochs 20 and 30, each evaluate 2,000
478 actual CPU bitstreams on 100 DIV2K centre crops across
479 D2/D4/D6 and released D12. On the same 99 images at 0.2
480 payload bpp, D6's RGB advantage over D2 changes from
481 0.183 to 0.107 dB; its epoch-30 difference from D4 is unre-
482 solved. Depth ordering is therefore not stable at these unfin-
483 ished checkpoints. The supplement reports both milestones.
484 Final same-recipe D12 and matched shared-exit training
485 remain necessary to isolate depth and weight sharing.

486 **The spatial cost of specialisation.** Patchwise expert selec-
487 tion also changes geometry and entropy boundaries. With
488 identical D6 weights, independently coded regions lose qual-
489 ity at matched rate; increasing context alone does not remove
490 their rate cost. In a separate fixed D2/D6 assignment, merg-
491 ing compatible neighbouring regions improves matched-rate
492 quality while retaining the same per-pixel depth map. These
493 controls, detailed in the supplement, make patchwise D12
494 and blind expert mixtures essential baselines before a selec-
495 tor receives credit. The eventual comparison must match
496 achieved rate and quality, charge signalling and grouping,
497 and measure the effect of expert occupancy on latency. It
498 will test when independent specialisation is worth more than
499 shared-prefix reuse.

500 6. Discussion

501 **Adaptation should pay for the decisions it changes.**
502 RegLIC identifies a budget-dependent role for spatial al-
503 location. Under source calibration, tight quality targets leave
504 room for useful placement. Near the shallowest-exit ceiling,
505 a blind mixture reaches nearly the same arithmetic saving. A

budget-conditioned switch between fixed, blind and learned
policies should therefore be chosen on validation data.

The remaining evidence gap. The source-calibrated study
measures available allocation value, while the fixed-control
replay leaves the routing increment unresolved. Content-
grouped calibration on untouched images is the next test
of that increment. Native bitstreams and paired wall-clock
measurements must then include entropy work, prediction,
tile movement, grouping and fallback. The current confi-
dence intervals cover sample variation at fixed weights, not
independent training seeds. Results are confined to intra
reconstruction: YCbCr-MSE and RGB diagnostics do not
establish perceptual equivalence, temporal consistency or
performance of DCVC-UF's chunk-based video path.

506 7. Conclusion

RegLIC makes synthesis depth a spatial decision inside
DCVC-UF. Aligned exits and a shrinking tile batch reuse a
common representation while omitting later block evalua-
tions. Most measured arithmetic saving comes from cheaper
depth; learned placement adds a smaller, budget-dependent
increment. Blind mixtures and fixed-control replay separate
these effects. Matched training and complete execution mea-
surements will determine when independent-depth experts
justify giving up shared computation.

References

- [1] Théophile Blard, Pierrick Philippe, Théo Ladune, Xiaoran Jiang, and Olivier Déforges. Spatial competition for low-complexity learned image compression. In *IEEE International Conference on Image Processing (ICIP)*, 2026. Accepted; author manuscript arXiv:2605.13243. 2
- [2] Mallesham Dasari, Kumara Kahatapitiya, Samir R. Das, Aruna Balasubramanian, and Dimitris Samaras. Swift: Adaptive video streaming with layered neural codecs. In *19th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, pages 103–118. USENIX Association, 2022. 2
- [3] Zhihao Hu and Dong Xu. Complexity-guided slimmable decoder for efficient deep video compression. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14358–14367, 2023. 2
- [4] Zhaoyang Jia, Bin Li, Jiahao Li, Wenxuan Xie, Linfeng Qi, Houqiang Li, and Yan Lu. Towards Practical Real-Time Neural Video Compression. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 12543–12552, 2025. 2
- [5] Jong Wook Kim, Suyong Bahk, TaeHwa Lee, HyunDong Cho, Donghyun Kim, Sung-Chang Lim, Jin Soo Choi, and Hui Yong Kim. Block-based learned image compression without blocking artifacts. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026. 2
- [6] Xiangtao Kong, Hengyuan Zhao, Yu Qiao, and Chao Dong. ClassSR: A General Framework to Accelerate Super-Resolution Networks by Data Characteristic. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 12016–12025, 2021. 1, 2
- [7] Jiahao Li, Wenxuan Xie, Zhaoyang Jia, Bin Li, Zongyu Guo, Xiaoyi Zhang, and Yan Lu. Ultra-Fast Neural Video Compression. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 41311–41321, 2026. 1, 2
- [8] Xintian Mao, Qingli Li, and Yan Wang. AdaRevD: Adaptive Patch Exiting Reversible Decoder Pushes the Limit of Image Deblurring. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 25681–25690, 2024. 2
- [9] Microsoft. DCVC-UF Image Training Recipe. Source repository, commit cbdae87a5445114cdc7f48816da63ea80bdeac40, 2026. 8
- [10] Calvin-Khang Ta, Praneet Singh, Tong Shao, and Peng Yin. MixCompress: Mixture of experts for variable rate learned image compression. In *European Conference on Computer Vision (ECCV)*, 2026. Accepted; author manuscript arXiv:2607.14334. 2
- [11] Lvfang Tao, Wei Gao, Ge Li, and Chenhao Zhang. Adanic: Towards practical neural image compression via dynamic transform routing. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 16879–16888, 2023. 1, 2
- [12] Shizun Wang, Jiaming Liu, Kaixin Chen, Xiaoqi Li, Ming Lu, and Yandong Guo. Adaptive Patch Exiting for Scalable Single Image Super-Resolution. In *European Conference on Computer Vision (ECCV)*, 2022. 1, 2
- [13] Qunliang Xing, Mai Xu, Tianyi Li, and Zhenyu Guan. Early exit or not: Resource-efficient blind quality enhancement for compressed images. In *European Conference on Computer Vision (ECCV)*, 2020. 2
- [14] Fei Yang, Luis Herranz, Yongmei Cheng, and Mikhail G. Mozerov. Slimmable Compressive Autoencoders for Practical Neural Image Compression. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021. 2
- [15] Yibo Yang and Stephan Mandt. Computationally-Efficient Neural Image Compression with Shallow Decoders. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023. 2
- [16] Yufeng Zhang, Wenrui Dai, Hang Yu, Shizhan Liu, Junhui Hou, Jianguo Li, and Weiyao Lin. ABC: Adaptive bayesnet structure learning for computational scalable multi-task image compression. arXiv:2506.15228, 2025. 2
- [17] Zhichen Zhang, Bolin Chen, Hongbin Lin, Jieliang Lin, Xu Wang, and Tiesong Zhao. ELFIC: A learning-based flexible image codec with rate-distortion-complexity optimization. In *Proceedings of the 31st ACM International Conference on Multimedia*, pages 9252–9261, 2023. 2