



THE PADDING POLICY CHANGES THE COMPUTE ESTIMATE

Four 256 cores · native image-pad16 / latent-pad4 versus CPU reference image-pad64



Layer-by-layer Conv2d trace · native-shaped means geometry only · no CUDA timing, stream or numerical parity claim