


PADDING CAN CONSUME THE DEPTH SAVING

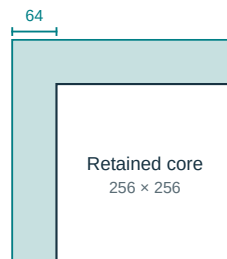
Bottom-right patch of a 2x2 grid · extra context need not mean extra padded area

a 32-pixel halo · 288 → 320



Same coded area: 320 × 320

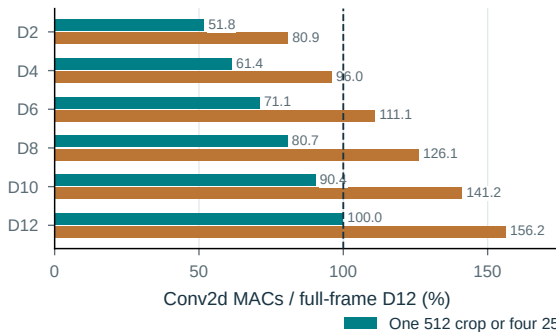
b 64-pixel halo · 320 → 320



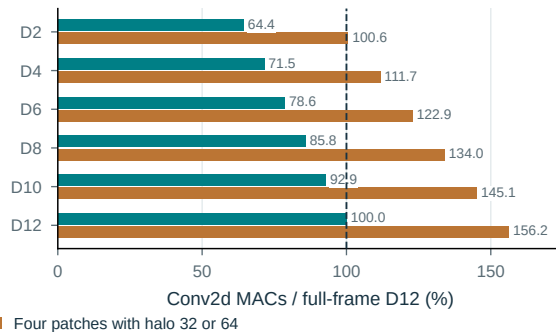
Same coded area: 320 × 320

Actual source context Replicated padding

c Neural decoder



d Encoder with reconstruction



Exact architecture/geometry accounting · no runtime, trained D8/D10 quality, or context-sufficiency claim