


a Selection and reporting use different reconstruction anchors and image support


SELECTION



Source / latent



e15 candidates +
released full reference




Padded source



Tile M (e15) + R (released)

Reference: released / 444-MSE

REPORTING



Source / latent



e15 mixed + full



Crop to original extent







Final mixed reconstruction

Reference: e15 / 444-MSE or 6:1:1


CPU audit: 255 shared encoder/entropy tensors unchanged; 143/143 inherited synthesis tensors changed.

Required comparison: reconstruct both anchors, use the same valid pixels and metric, then verify the selected output.




Source-informed search Router Bayer dither Uniform depth

a Compare outputs under the same loss cap




265 pairs at every cap · fine-tuned full-depth reference · retrospective selection with explicit fallback

b The routing margin survives this control



a Independent codecs: depth changes the synthesis network, training changes the whole codec



D2 / D4 / D6: pinned official recipe, all codec weights learned jointly

A same-recipe D12 control is still needed to separate depth from training provenance.

b Proposed deployment: choose before encoding, group work by expert, transmit the choice






The decoder reads the expert ID and uses that expert's entropy model and synthesis transform.

Measure: full bytes · encoder + decoder wall time · final cropped quality · bank memory




Deployment is a proposed design. The glyphs and queues are schematic; they do not show measured routes or throughput.

DEPTH-STUDY DESIGN SPACE

Exact parameter counts | Analytical runtime scenarios — not measurements or fitted forecasts



f is hypothetical trunk time, h is extra overhead. Dashed curves have no PSNR, bitrate or empirical confidence interval.



265 frame-QP pairs · equal frame weighting · 95% sequence-cluster intervals

Archived uniform reconstructions on padded YCbCr 4:4:4; these are shared exits, not independently trained shallow codecs.