

a Selection and reporting use different reconstruction anchors and image support


SELECTION



Source / latent



e15 candidates +
released full reference



Padded source



Tile M (e15) + R (released)

Reference: released / RGB

REPORTING



Source / latent



e15 mixed + full



Crop to original extent







Final mixed reconstruction


Reference: e15 full / RGB or YUV

CPU audit: 255 shared encoder/entropy tensors unchanged; 143/143 inherited synthesis tensors changed.

Required comparison: reconstruct both anchors, use the same valid pixels and metric, then verify the selected output.



a Independent codecs: depth changes the synthesis network, training changes the whole codec



D2 / D4 / D6: pinned official recipe, all codec weights learned jointly

A same-recipe D12 control is still needed to separate depth from training provenance.

b Proposed deployment: choose before encoding, group work by expert, transmit the choice



The decoder reads the expert ID and uses that expert's entropy model and synthesis transform.

Measure: full bytes · encoder + decoder wall time · final cropped quality · bank memory

Deployment is a proposed design. The glyphs and queues are schematic; they do not show measured routes or throughput.