



a A shared representation, spatially variable synthesis



b Only surviving tiles traverse the next block pair







Frozen encoder + entropy model


The synthesis network and exit adapters are fine-tuned.

Shared head couples neighbouring tiles


Evaluate the final mixed reconstruction after cropping.




a Achieved quality differs across rules




b Reporting differs from selection




c Losses above the nominal target



a The source is available at the encoder, but the error table must be computed



b Recorded table construction time



Historical timing: 8 frames $\times$ 3 repetitions at QP32. No new end-to-end latency measurement.

c Calibration changes the cost

Fixed, held-out β


Router uses decoded features.
No per-frame source table is required.

Per-frame source-calibrated β


The plotted budget sweep uses $M[t, k]$.
A fast router does not remove that cost.

a Spatial decisions converge as the quality budget relaxes

videoSRC05_1920x1080_25.yuv · QP32 · archived source luma



a Shared early exit: one representation, nested synthesis



b Independent depth controls: each codec learns its own representation



Same-recipe D12 is an additional control. The released anchor alone does not isolate depth.

c Later routing study: select before encoding; signal which expert owns the payload

